

JOHN DOE

AI-DRIVEN FULL-STACK DEVELOPER / AI ENGINEER

Full-stack Engineering / AI Product Integration / Backend APIs / React / Next.js / Python / FastAPI / Product & Delivery

 July 4, 1995  San Francisco, CA  +(84) ...  john.doe@email.com  linkedin.com/in/hanhtran47  github.com/hanh-trannguyen

JOHN DOE'S CV

- Email: john.doe@email.com
- Location: San Francisco, CA
- Website: rendercv.com
- LinkedIn: rendercv
- GitHub: rendercv

WELCOME TO RENDERCV

RenderCV reads a CV written in a
YAML file, and generates a PDF
with professional typography.
Each section title is arbitrary.
You can choose any of the 9 entry
types for each section.
Markdown syntax is supported
everywhere. This is **bold**, *italic*,
and link.

EDUCATION

Princeton University, Computer Science

PhD

Princeton, NJ

Sept 2018 – May 2023

- Thesis: Efficient Neural
Architecture Search for Resource-
Constrained Deployment
- Advisor: Prof. Sanjeev Arora
- NSF Graduate Research
Fellowship, Siebel Scholar (Class
of 2022)

Boğaziçi University, Computer Engineering

BS

Istanbul, Türkiye

Sept 2014 – June 2018

- GPA: 3.97/4.00, Valedictorian
- Fulbright Scholarship recipient for
Graduate Studies

EXPERIENCE

Nexus AI, Co-Founder & CTO

San Francisco, CA

June 2023 – present

3 years 1 month

- Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime
- Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund
- Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions
- Developed proprietary inference optimization reducing latency by 73% compared to baseline

NVIDIA Research, Research Intern

Santa Clara, CA

May 2022 – Aug 2022

4 months

- Designed sparse attention mechanism reducing transformer memory footprint by 4.2x
- Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)

Google DeepMind, Research Intern

London, UK

May 2021 – Aug 2021

4 months

- Developed reinforcement learning algorithms for multi-agent coordination
- Published research at top-tier venues with significant academic impact
- ICML 2022 main conference paper, cited 340+ times within two years
- NeurIPS 2022 workshop paper on emergent communication protocols
- Invited journal extension in JMLR (2023)

Apple ML Research, Research Intern

Cupertino, CA

May 2020 – Aug 2020

4 months

- Created on-device neural network compression pipeline deployed across 50M+ devices
- Filed 2 patents on efficient model quantization techniques for edge inference

Microsoft Research, Research Intern

Redmond, WA

May 2019 – Aug 2019

4 months

- Implemented novel self-supervised learning framework for low-resource language modeling
- Research integrated into Azure Cognitive Services, reducing training data requirements by 60%

PROJECTS

FlashInfer

Jan 2023 – present

Open-source library for high-performance LLM inference kernels

- Achieved 2.8x speedup over baseline attention implementations on A100 GPUs
- Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors

NeuralPrune

Jan 2021

Automated neural network pruning toolkit with differentiable masks

- Reduced model size by 90% with less than 1% accuracy degradation on ImageNet
- Featured in PyTorch ecosystem tools, 4,200+ GitHub stars

PUBLICATIONS

Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models

July 2023

John Doe, Sarah Williams, David Park

10.1234/neurips.2023.1234
(NeurIPS 2023)

Neural Architecture
Search via
Differentiable
Pruning

Dec 2022
James Liu, *John Doe*
10.1234/neurips.2022.5678
(NeurIPS 2022, Spotlight)

Multi-Agent
Reinforcement
Learning with
Emergent

Communication
July 2022
Maria Garcia, *John Doe*, Tom
Anderson
10.1234/icml.2022.9012 (ICML
2022)

On-Device Model
Compression via
Learned Quantization

May 2021
John Doe, Kevin Wu
10.1234/iclr.2021.3456 (ICLR 2021,
Best Paper Award)

SELECTED
HONORS

- MIT Technology Review 35 Under 35 Innovators (2024)
- Forbes 30 Under 30 in Enterprise Technology (2024)
- ACM Doctoral Dissertation Award Honorable Mention (2023)
- Google PhD Fellowship in Machine Learning (2020 – 2023)
- Fulbright Scholarship for Graduate Studies (2018)

SKILLS

Languages: Python, C++, CUDA,
Rust, Julia
ML Frameworks: PyTorch, JAX,
TensorFlow, Triton, ONNX
Infrastructure: Kubernetes, Ray,
distributed training, AWS, GCP

Research Areas: Neural
architecture search, model
compression, efficient inference,
multi-agent RL

PATENTS

1. Adaptive Quantization for
Neural Network Inference
on Edge Devices (US
Patent 11,234,567)
2. Dynamic Sparsity Patterns
for Efficient Transformer
Attention (US Patent
11,345,678)
3. Hardware-Aware Neural
Architecture Search
Method (US Patent
11,456,789)

INVITED TALKS

1. Scaling Laws for Efficient
Inference — Stanford HAI
Symposium (2024)
2. Building AI Infrastructure
for the Next Decade —
TechCrunch Disrupt (2024)
3. From Research to
Production: Lessons in ML
Systems — NeurIPS
Workshop (2023)
4. Efficient Deep Learning: A
Practitioner's Perspective
— Google Tech Talk (2022)