

EDUCATION

Princeton University

Boğaziçi University

SKILLS

LANGUAGES

Python C++ CUDA Rust Julia

ML FRAMEWORKS

PyTorch JAX TensorFlow Triton

ONNX

INFRASTRUCTURE

Kubernetes Ray distributed training

AWS GCP

RESEARCH AREAS

Neural architecture search

model compression efficient inference

multi-agent RL

JOHN DOE

Email: john.doe@email.com

Location: San Francisco, CA

WELCOME TO RENDERCV

RenderCV reads a CV written in a YAML file, and generates a PDF with professional typography.

Each section title is arbitrary.

You can choose any of the 9 entry types for each section.

Markdown syntax is supported everywhere. This is **bold**, *italic*, and [link](#).

EXPERIENCE

Nexus AI

San Francisco, CA · Co-Founder & CTO

- Built foundation model infrastructure serving 2M+ monthly API requests with 99.97% uptime
- Raised \$18M Series A led by Sequoia Capital, with participation from a16z and Founders Fund
- Scaled engineering team from 3 to 28 across ML research, platform, and applied AI divisions
- Developed proprietary inference optimization reducing latency by 73% compared to baseline

NVIDIA Research

Santa Clara, CA · Research Intern

- Designed sparse attention mechanism reducing transformer memory footprint by 4.2x
- Co-authored paper accepted at NeurIPS 2022 (spotlight presentation, top 5% of submissions)

Google DeepMind

London, UK · Research Intern

- Developed reinforcement learning algorithms for multi-agent coordination
- Published research at top-tier venues with significant academic impact - ICML 2022 main conference paper, cited 340+ times within two years - NeurIPS 2022 workshop paper on emergent communication protocols - Invited journal extension in JMLR (2023)

Apple ML Research

Cupertino, CA · Research Intern

- Created on-device neural network compression pipeline deployed across 50M+ devices
- Filed 2 patents on efficient model quantization techniques for edge inference

Microsoft Research

Redmond, WA · Research Intern

- Implemented novel self-supervised learning framework for low-resource language modeling
- Research integrated into Azure Cognitive Services, reducing training data requirements by 60%

PROJECTS

FlashInfer

Open-source library for high-performance LLM inference kernels

- Achieved 2.8x speedup over baseline attention implementations on A100 GPUs
- Adopted by 3 major AI labs, 8,500+ GitHub stars, 200+ contributors

NeuralPrune

Automated neural network pruning toolkit with differentiable masks

- Reduced model size by 90% with less than 1% accuracy degradation on ImageNet
- Featured in PyTorch ecosystem tools, 4,200+ GitHub stars

PUBLICATIONS

Sparse Mixture-of-Experts at Scale: Efficient Routing for Trillion-Parameter Models

John Doe, Sarah Williams, David Park · NeurIPS 2023

Neural Architecture Search via Differentiable Pruning

James Liu, John Doe · NeurIPS 2022, Spotlight

Multi-Agent Reinforcement Learning with Emergent Communication

Maria Garcia, John Doe, Tom Anderson · ICML 2022

On-Device Model Compression via Learned Quantization

John Doe, Kevin Wu · ICLR 2021, Best Paper Award

 SELECTED HONORS

MIT Technology Review 35 Under 35 Innovators (2024)

Forbes 30 Under 30 in Enterprise Technology (2024)

ACM Doctoral Dissertation Award Honorable Mention (2023)

Google PhD Fellowship in Machine Learning (2020 – 2023)

Fulbright Scholarship for Graduate Studies (2018)

 PATENTS

Adaptive Quantization for Neural Network Inference on Edge Devices (US Patent 11,234,567)

Dynamic Sparsity Patterns for Efficient Transformer Attention (US Patent 11,345,678)

Hardware-Aware Neural Architecture Search Method (US Patent 11,456,789)

 INVITED TALKS

Scaling Laws for Efficient Inference — Stanford HAI Symposium (2024)

Building AI Infrastructure for the Next Decade — TechCrunch Disrupt (2024)

From Research to Production: Lessons in ML Systems — NeurIPS Workshop (2023)

Efficient Deep Learning: A Practitioner's Perspective — Google Tech Talk (2022)